

RangeQueue: A Priority Queue algorithm

Abhay Shankar K: cs21btech11001 Kartheek Tamma: cs21btech11028
Rajapu Jayachandra Naidu: cs21btech11050

December 3, 2023

Abstract

This document is a report for the Final Project of the course CS5300: Parallel and Concurrent Programming, regarding a novel implementation of a lock-free data structure.

Introduction and Related Work

A priority queue supports two operations:

- **insert(key)**: Insert a key into the queue.
- **deleteMin()**: Delete and return the key with the lowest key from the queue.

Our implementation is based on the TSLQueue data structure proposed by Rukundo and Tsigas. In a TSLQueue, keys are stored in a binary search tree, and each node contains a **next** pointer to the node with the next lowest key. This allows for **insert** in log time, and constant time **deleteMin**, while still being parallelizable.

While the TSLQueue only stores a single key in each node, we generalize it to hold a range of keys per node, using a bitvector

Optimisations

The RangeQueue is an improvement over the TSLQueue in the following ways:

- Cache locality: Since each node represents a range of keys, the number of nodes in the queue is reduced. This reduces the number of cache misses.
- Fewer CAS operations:
 - The TSLQueue requires 2 CAS operations per insertion, because each insertion modifies a linked list and a tree. RangeQueue requires either 1 CAS operation, or 1 CAS and 1 **FETCH_OR** per insertion. This will be elaborated upon in the Algorithm section.
 - The TSLQueue requires 3 CAS operations per deletion, one to logically delete the node by marking a bit, and two to modify pointers. RangeQueue again requires fewer CAS's because not every deletion requires pointer rearrangement.
- Lesser memory footprint: We use bitvectors instead of storing the actual value.

Overview

A threaded tree is a tree in which every node has a pointer to its successor in an inorder traversal. Such a concept can be generalised from key-based nodes to range-based nodes trivially, as non-overlapping ranges have a total ordering.

Each node in the RangeQueue contains

- **base**: The value of the the least key that could be inserted into it. Must be divisible by **BUF_SIZE**.
- **next**: A pointer to the next node in the queue.
- **parent, left, right**: Pointers to the parent, left child, and right child of the node respectively. Standard BST terminology.
- **buf**: A bitvector of size **BUF_SIZE** that indicates whether the key at an index has been inserted.

- **ins**: A flag indicating whether the node is currently undergoing insertion. Identical to TSLQueue.

The RangeQueue has a **head** pointer that points to the first node in the queue. The **head** pointer is a dummy node that has a key of -1 and a **next** pointer to the first node in the queue. The **head** pointer is used to simplify the implementation of the **deleteMin()** operation and address the ABA problem.

The RangeQueue also has a dummy root. The root has a **next** pointer to the actual root node of the tree.

Both the dummy nodes are used in TSLQueue, for ease of proof of correctness.

Algorithm

Insertion

- A thread descends the BST to find the node that contains the key to be inserted.
 - If such a node exists, it atomically sets the corresponding bit in the **buf** bitvector. If the bit is already set, returns duplication error.
 - * If a node is logically deleted (i.e. **buf** is all zeroes), it's existence will not be reported, i.e. **InsertSearch** returns **existingNode = null**.
 - Otherwise, it allocates a new node with the appropriate bit set, and inserts it into the RangeQueue. This is identical to the TSLQueue insert, with an similar **Seek** struct.
 - * Set **next** to next node in the linked list.
 - * Atomically insert into the linked list by setting **next** pointer of the previous node to the new node.
 - * Set **parent** pointer of the new node to the parent node.
 - * Atomically insert into the BST by setting appropriate **left** or **right** pointer of the parent node to the new node.

Deletion

- Deletion uses the **head** pointer to access the minimal node, and resets the highest set bit in the **buf** bitvector.
 - This is a single atomic operation - upon failure, the thread simply retries delete.
 - If after deletion, the **buf** bitvector is all zeroes, the node is said to be logically deleted, unlike in TSLQueue where there is a specific tag on the **next** pointer to indicate the same. With some nonzero probability, the thread physically deletes the node. Physical deletion is identical to TSLQueue.
 - * Atomically delete from the linked list by setting **next** pointer of **head** to the next node.
 - * Atomically delete from the BST by setting **left** pointer of the parent node to the next node.

References

Rukundo, A., Tsigas, P. (2021). TSLQueue: An Efficient Lock-Free Design for Priority Queues. In: Sousa, L., Roma, N., Tomás, P. (eds) Euro-Par 2021: Parallel Processing. Euro-Par 2021. Lecture Notes in Computer Science(), vol 12820. Springer, Cham. https://doi.org/10.1007/978-3-030-85665-6_24